

CS777 Term Paper Presentation

Exploration of K8s

Venkata Sandeep Yerra
Boston University


Evolution, Architecture and MLOps

Agenda

- ❖ Introduction
- ❖ Historical Background
- ❖ Evolution of Containers and Rise of Kubernetes
- ❖ Kubernetes Architecture
- ❖ Demo: MLOps
- ❖ Demo: Monitoring

Introduction

- ❖ What is a container?
 - ❖ Lightweight, portable box that has everything
 - ❖ App itself along with code
 - ❖ Libraries
 - ❖ Settings
- ❖ Why do we need a container
 - ❖ Differences in Operating systems and underlying infrastructure are abstracted
 - ❖ Much more light weight compared to Virtual machines which also ship operating system
- ❖ What is Kubernetes
 - ❖ Framework to run distributed systems resiliently
- ❖ Why do we Kubernetes
 - ❖ Automates operational tasks of container management, deployment of applications, rolling changes, scaling, monitoring



Historical Background of Kubernetes

- ❖ The story of Kubernetes begins with its roots in a project at Google. Let's take a quick journey through its history
- ❖ **2003-2014: The Google Influence**
 - ❖ The concept underlying Kubernetes is largely based on Google's internal platform called Borg, which revolutionized the way software was deployed and managed at scale within Google since around 2003
 - ❖ Borg allowed Google to efficiently run containers – which are like lightweight, standalone packages of software – across their massive server fleets
 - ❖ **Omega** was a ground up software solution developed internally at google to be more flexible than Borg
- ❖ **2014: Birth of Kubernetes**
 - ❖ In mid-2014, Google decided to build an open-source version of their internal tools, leading to the birth of Kubernetes
 - ❖ Google partnered with the Linux Foundation to form the Cloud Native Computing Foundation (CNCF) to govern Kubernetes independently.
- ❖ **2015-Present: Growth and Adoption**
 - ❖ Kubernetes quickly gained popularity in the tech community for its efficiency in managing containerized applications
 - ❖ It became synonymous with microservices architectures and cloud-native technologies
- ❖ Today, Kubernetes is widely adopted by numerous companies, from startups to large enterprises, for orchestrating containers in production environment

Examples

Dockerfile

```
FROM python:3.8-slim-bullseye

ARG WORKDIR=/mlflow
RUN mkdir /mlflow
WORKDIR ${WORKDIR}

ENV LC_ALL=C.UTF-8
ENV LANG=C.UTF-8

RUN echo "export LC_ALL=$LC_ALL" >> /etc/profile.d/locale.sh
RUN echo "export LANG=$LANG" >> /etc/profile.d/locale.sh

COPY requirements.txt ${WORKDIR}
RUN pip install -U pip && \
    pip install --no-cache-dir -r requirements.txt

EXPOSE 5000

ENV DB_NAME=postgres
ENV DB_USERNAME=postgres
ENV DB_HOST=127.0.0.1
ENV DB_PASSWORD=password
ENV DEFAULT_ARTIFACT_ROOT=gs://example

ENTRYPOINT mlflow server \
    --host=0.0.0.0 \
    --port=5000 \
    --backend-store-uri=postgresql://${DB_USERNAME}:${DB_PASSWORD}@${DB_HOST}:5432/${DB_NAME} \
    --default-artifact-root=${DEFAULT_ARTIFACT_ROOT}
```

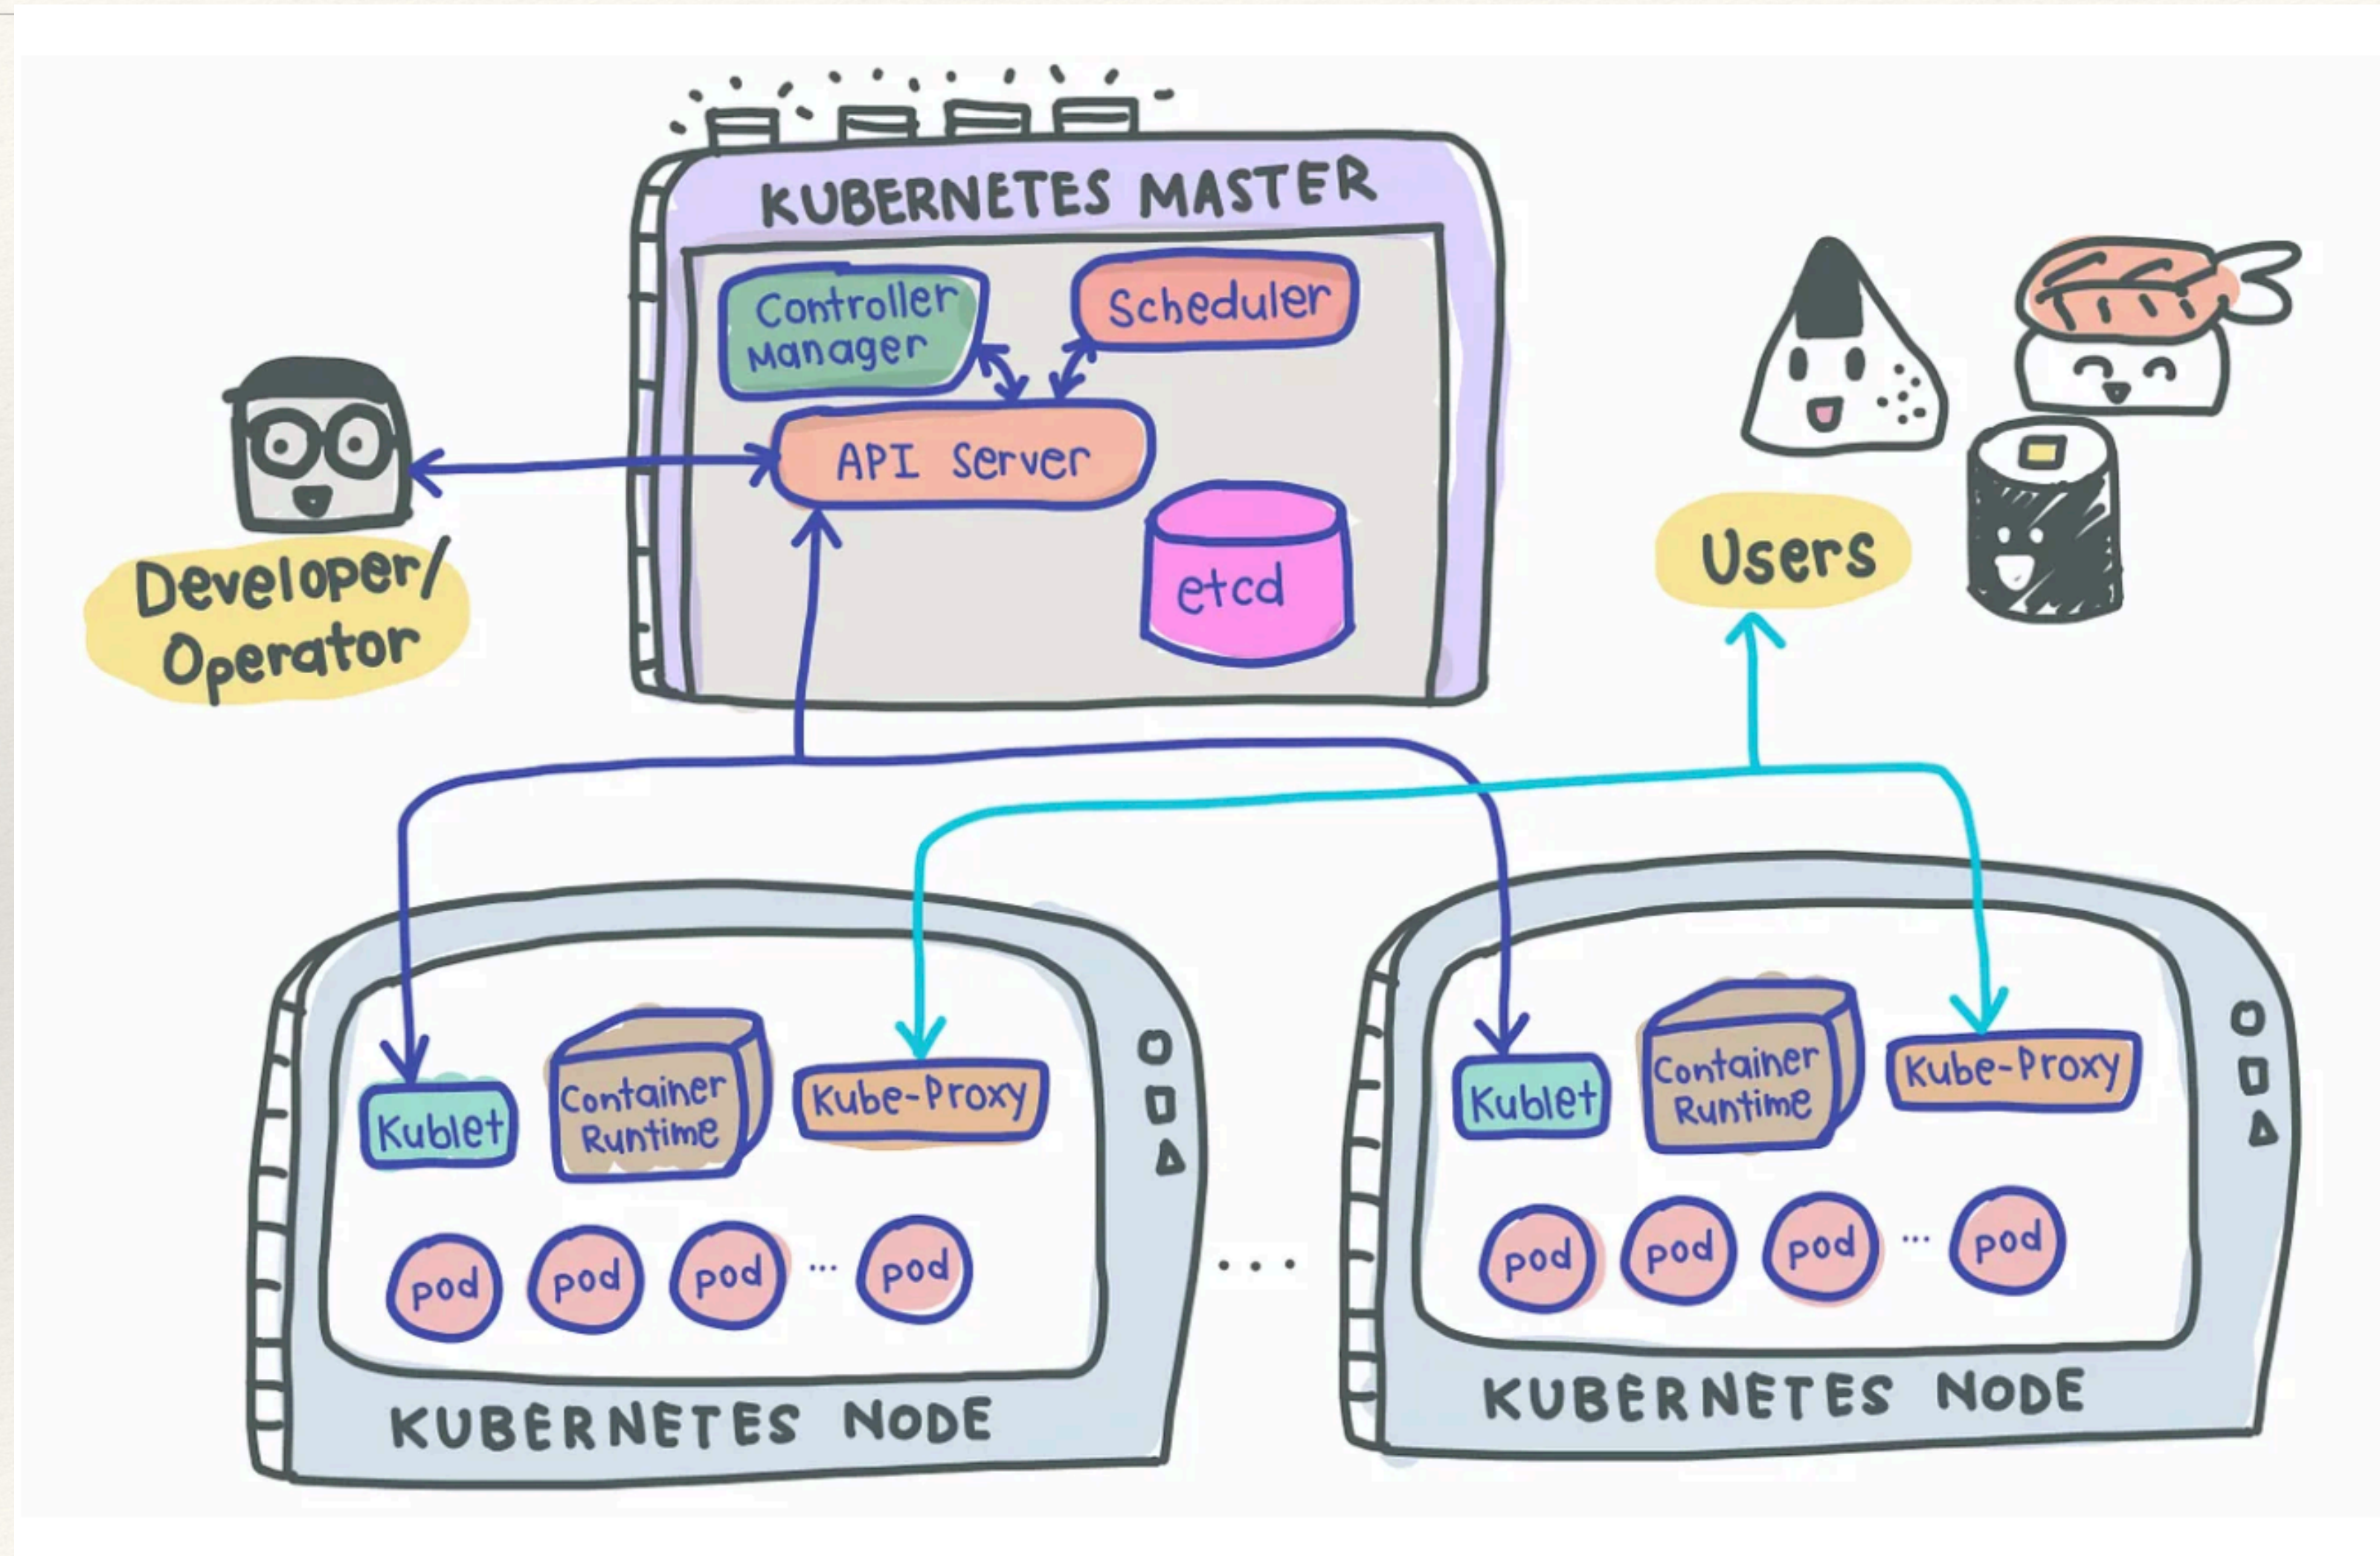
Kubernetes Deployment File

```
deployment_mlflow.yaml — Edit
# deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mlflow-tracking-server
  labels:
    app: mlflow-tracking-server
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mlflow-tracking-server-pods
  # Pod configurations defined here in `template`
  template:
    metadata:
      labels:
        app: mlflow-tracking-server-pods
    spec:
      containers:
        - name: mlflow-tracking-server-pod
          image: sandeepyerra/my-image:v2.1
          ports:
            - containerPort: 5000
          resources:
            limits:
              memory: 1Gi
              cpu: "2"
            requests:
              memory: 1Gi
              cpu: "1"
          imagePullPolicy: Always
          env:
            - name: DB_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mlflow-postgresql-credentials
                  key: postgresql-password
            - name: DB_USERNAME
              valueFrom:
                configMapKeyRef:
                  name: mlflow-configmap
                  key: DB_USERNAME
```

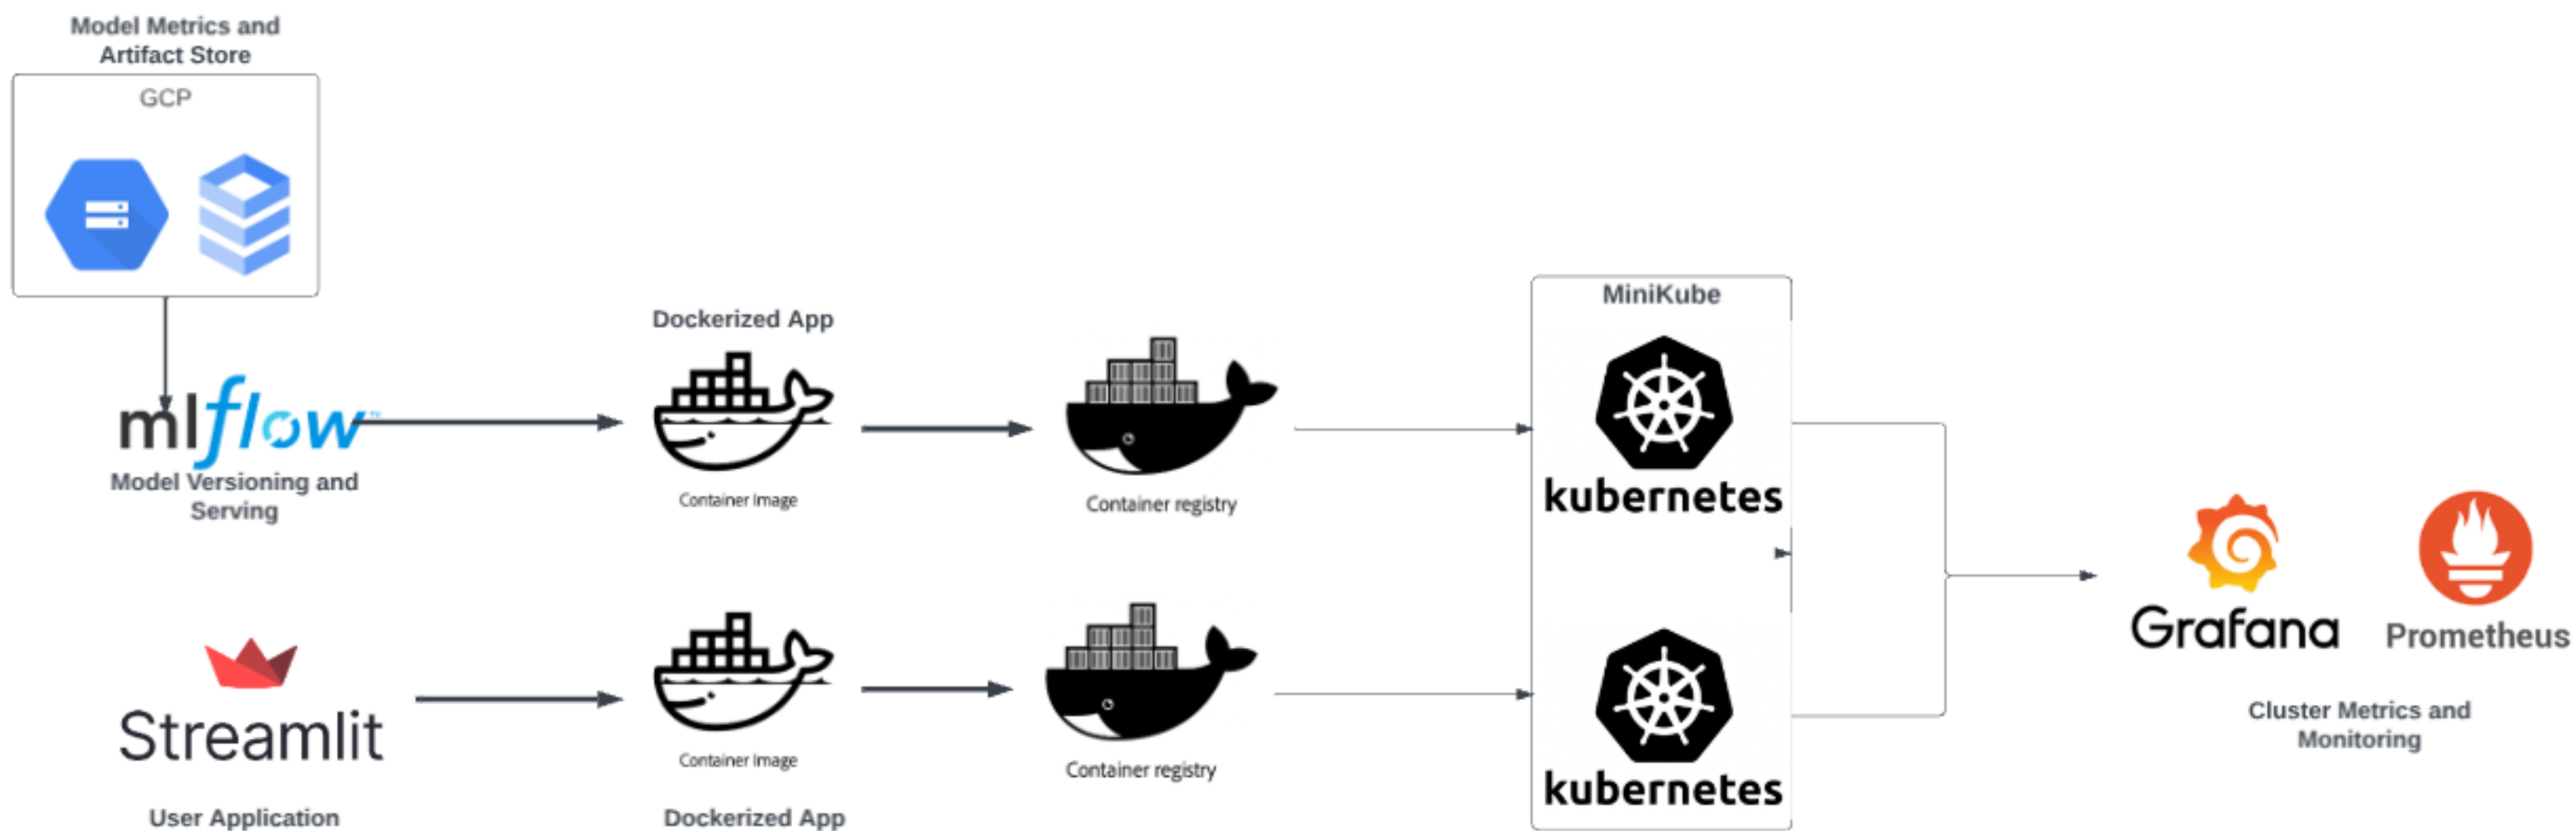

Kubernetes Architecture

Key Components

- Kubernetes Master
 - API Server
 - Controller Manager
 - Scheduler
 - ETCD
- Kubernetes Node
 - Kubelet
 - Container runtime
 - Kube-proxy
 - Pod



Demo Architecture



Demo Flow

- ❖ Flow of Demo
 - ❖ Kubernetes Features
 - ❖ Self-healing
 - ❖ Scaling
 - ❖ Model Development and Experiment Tracking
 - ❖ Model serving
 - ❖ Cluster Metrics Monitoring using Grafana and Prometheus

Thank You

